


Efficient Code Cache Management for Dynamic Multi-Tiered Compilation Systems

Tobias Hartmann, ETH Zurich, Oracle Corp.

Albert Noll, Oracle Corporation

Thomas R. Gross, ETH Zurich

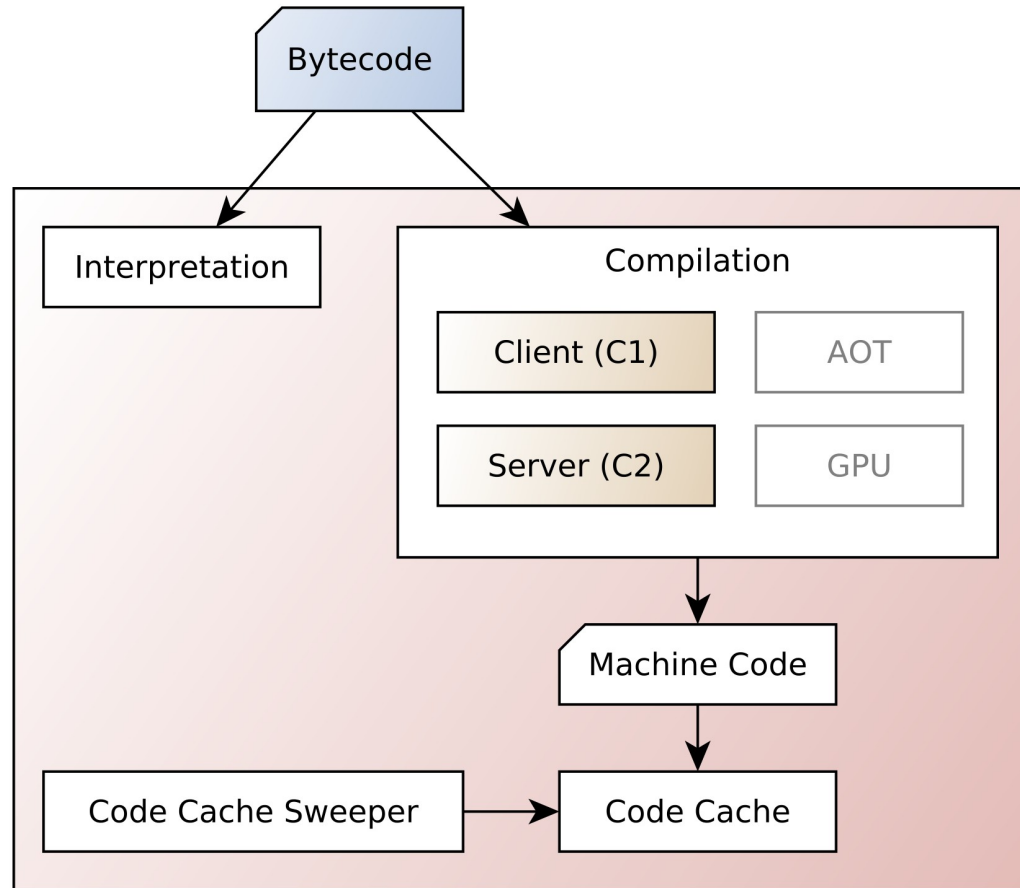
Introduction



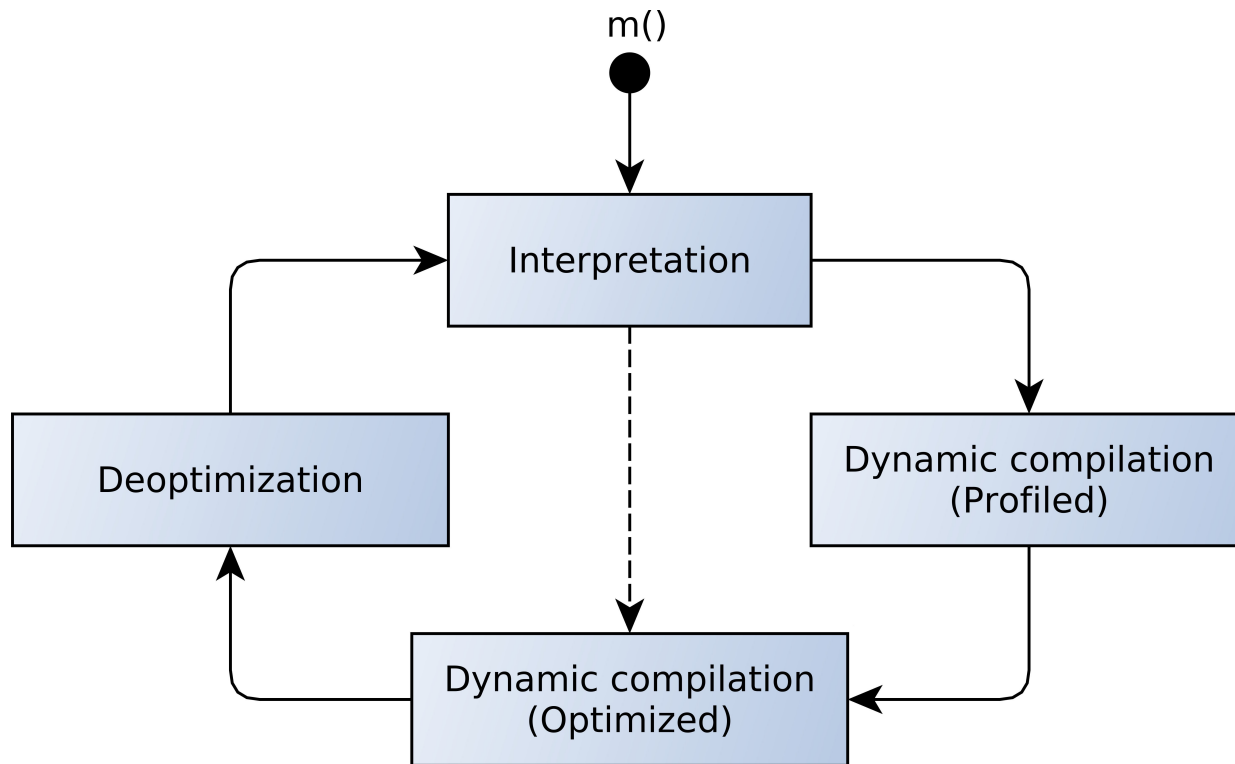
Outline

- Hotspot™ JVM
- Design
- Implementation
- Evaluation
- Conclusion

Hotspot™ JVM




Dynamic compilation in the JVM




History


■ JDK 6



■ JDK 7 / 8

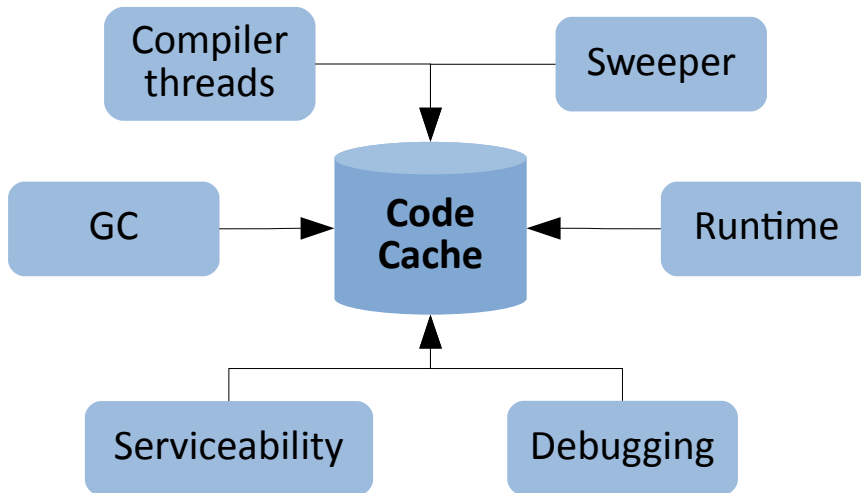


■ JDK 9 / Future



Code cache


■ Central component




■ Continuous chunk of memory

- Fixed size
- Bump pointer allocation with free list

Challenges

- With tiered compilation amount of code increased by **2-4 X**
- All code in one cache
 - **Different types** and characteristics
 - Access to specific code: **full iteration** 
- Code cache fragmentation

Challenges

- With tiered compilation amount of code increased by **2-4 X**
- All code in one cache
 - **Different types** and characteristics
 - Access to specific code: **full iteration** 
- Code cache fragmentation
- Solution: **Segmented Code Cache**


Properties of compiled code

- **Lifetime**
- **Size**
- **Cost of generation**

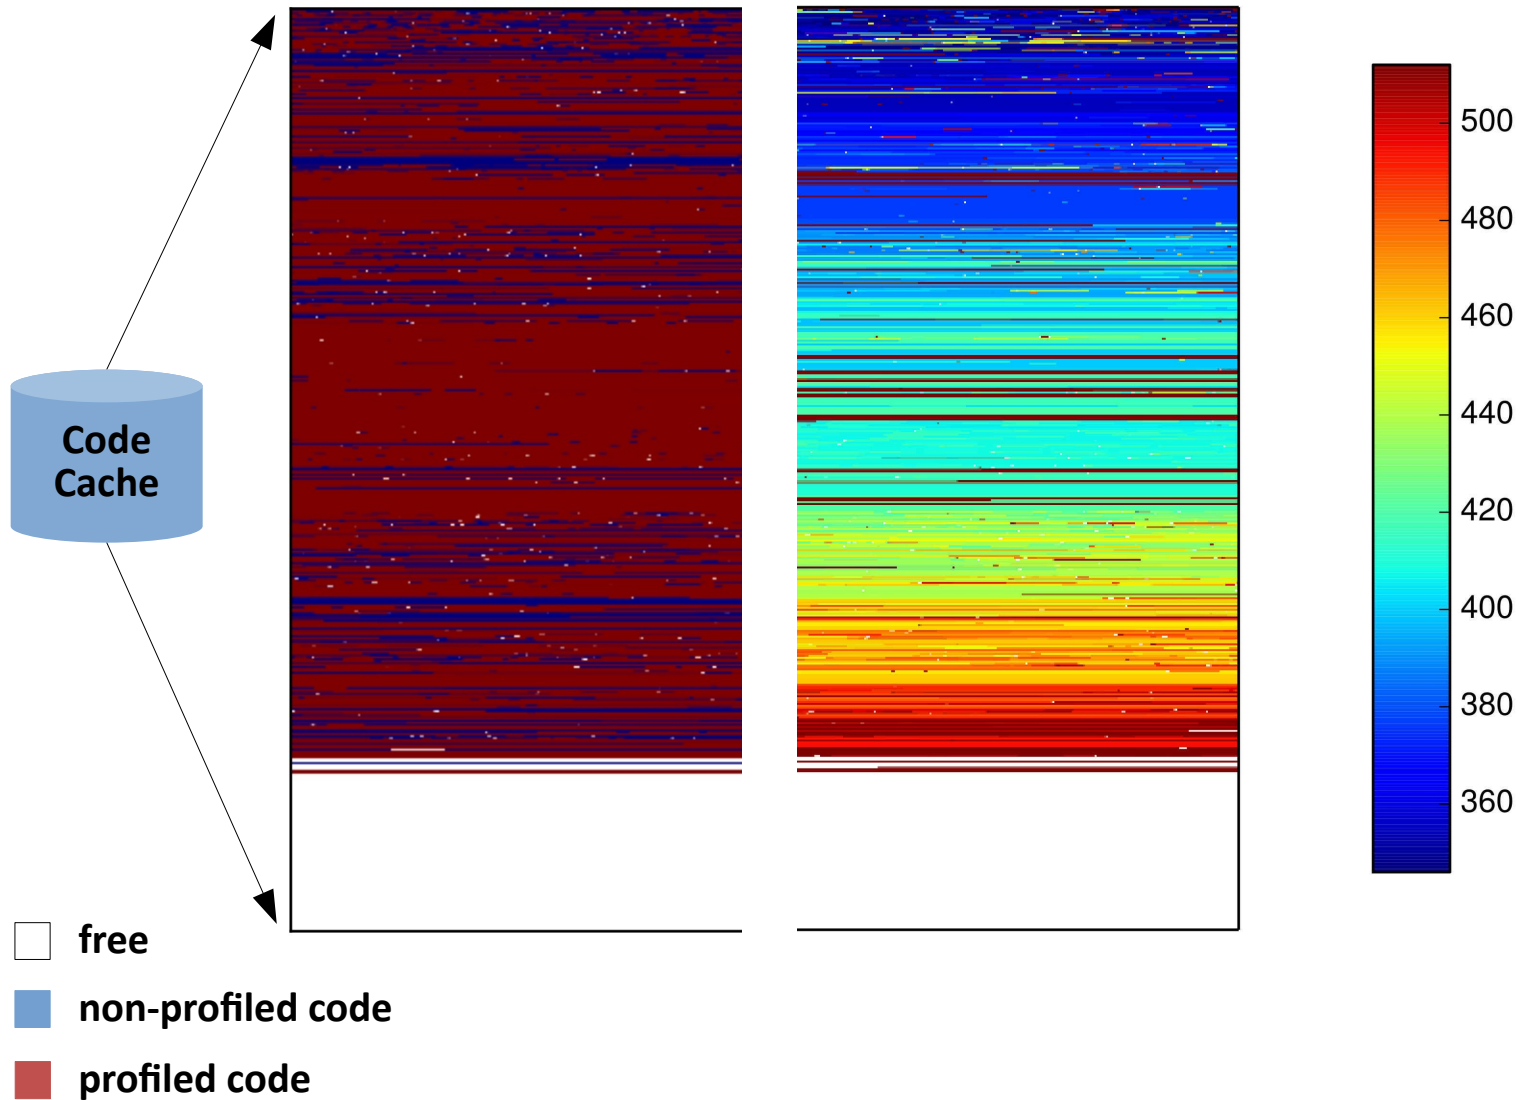
Types of compiled code

- **Non-method code**
- **Profiled method code**
 - Instrumented (C1)
 - **Limited lifetime**
- **Non-profiled method code**
 - Highly optimized (C2)
 - **Long lifetime**

Code cache fragmentation

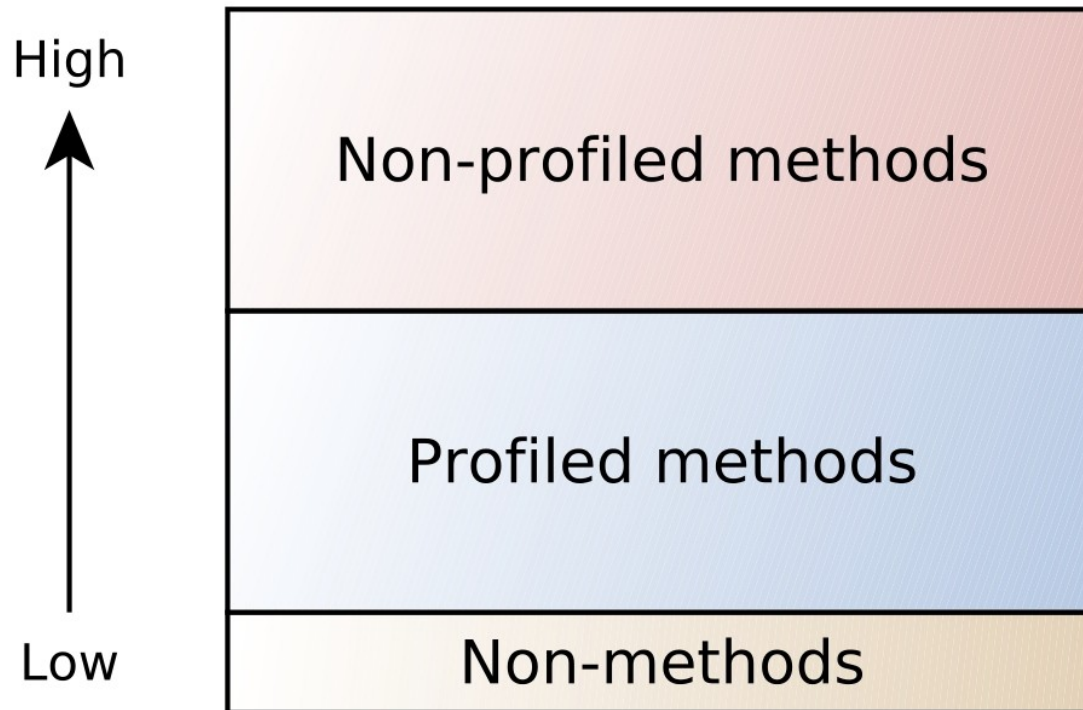


Hotness of code



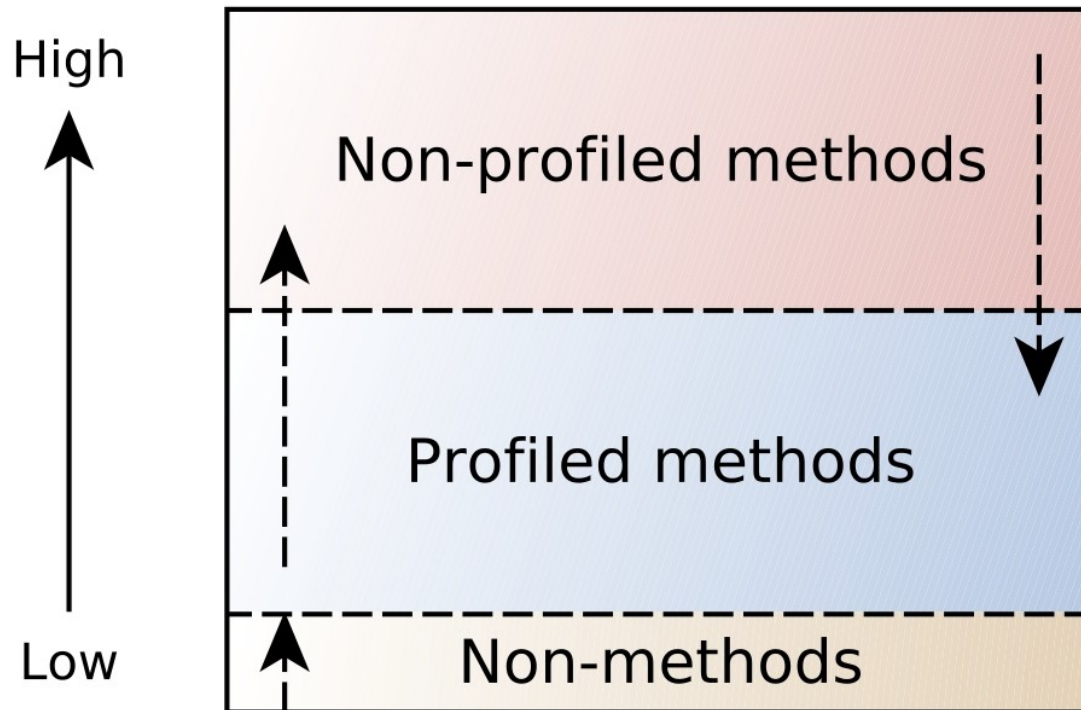
Segmented code cache

- Dividing code cache into distinct **segments**



Dynamic resizing

- Allowing the segments to **resize**



Implementation

- **Two prototype implementations**



- Fully functional
- With and without **resizing**

- **Corner cases**

- Small code cache sizes
- Different compiler configurations
- Code cache sweeper

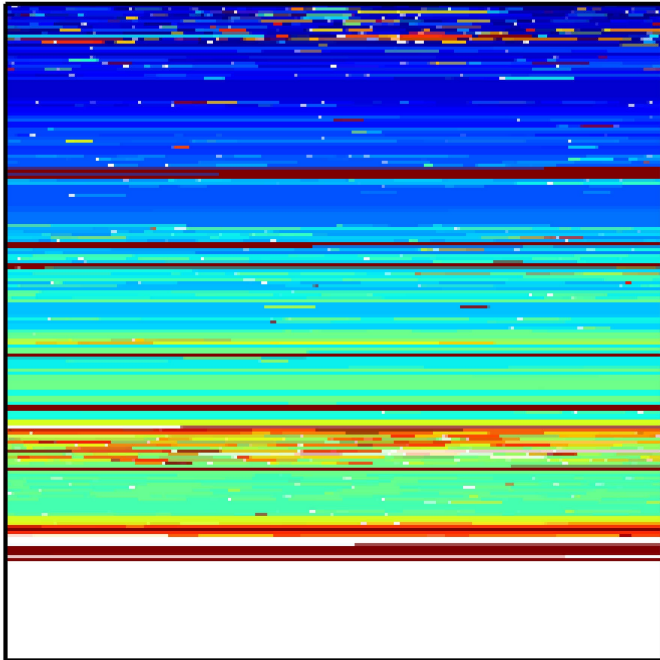
- **Several **optimizations** possible**

Code cache fragmentation




- ☐ free
- ☒ non-profiled code
- ☒ profiled code


Hotness of code



profiled code



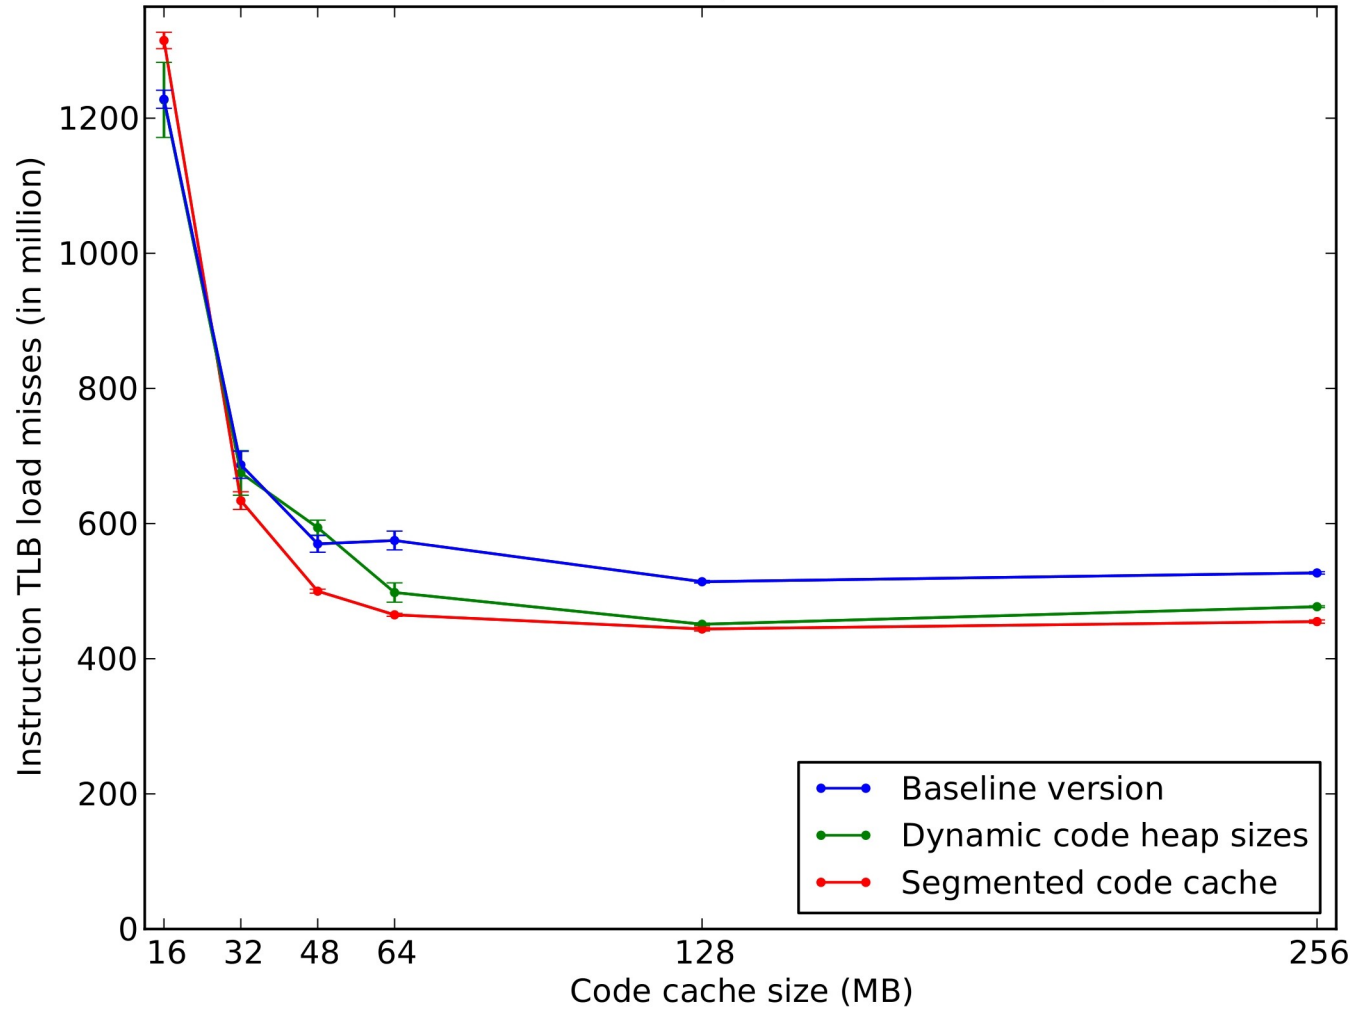
non-profiled code



Performance evaluation


- Segmented code cache
- Segmented code cache with dynamic resizing
- Hardware setup
 - 4 Intel Xeon E7-4830 CPUs at 2.13 GHz with 24 MB cache
 - 64 GB main memory

Instruction TLB




Instruction TLB

- 19 %

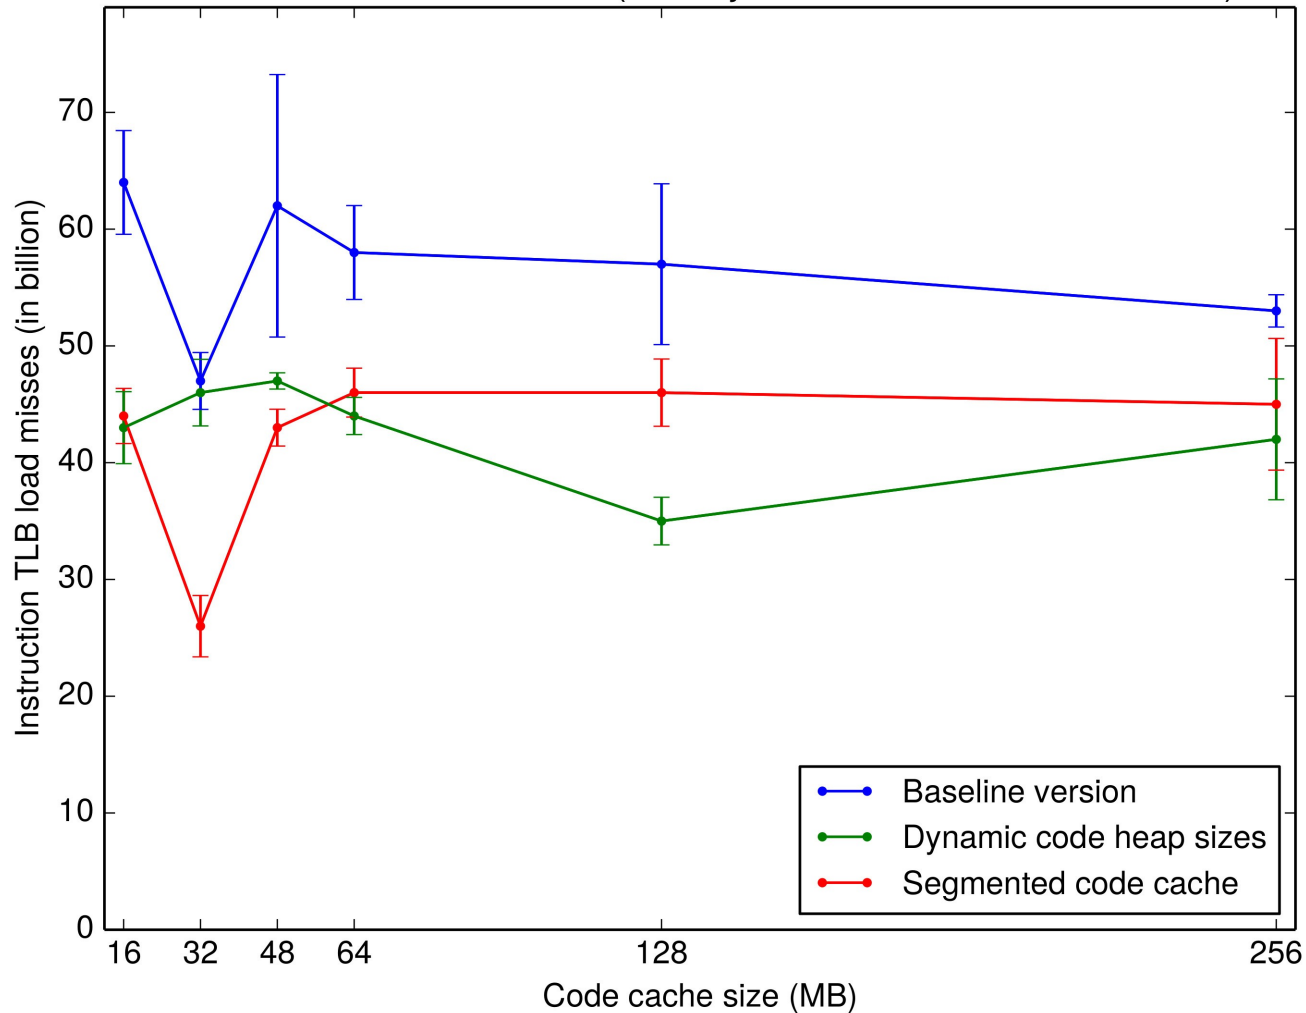


Instruction TLB (long running)




Instruction TLB (long running)

- 44 %




Instruction cache

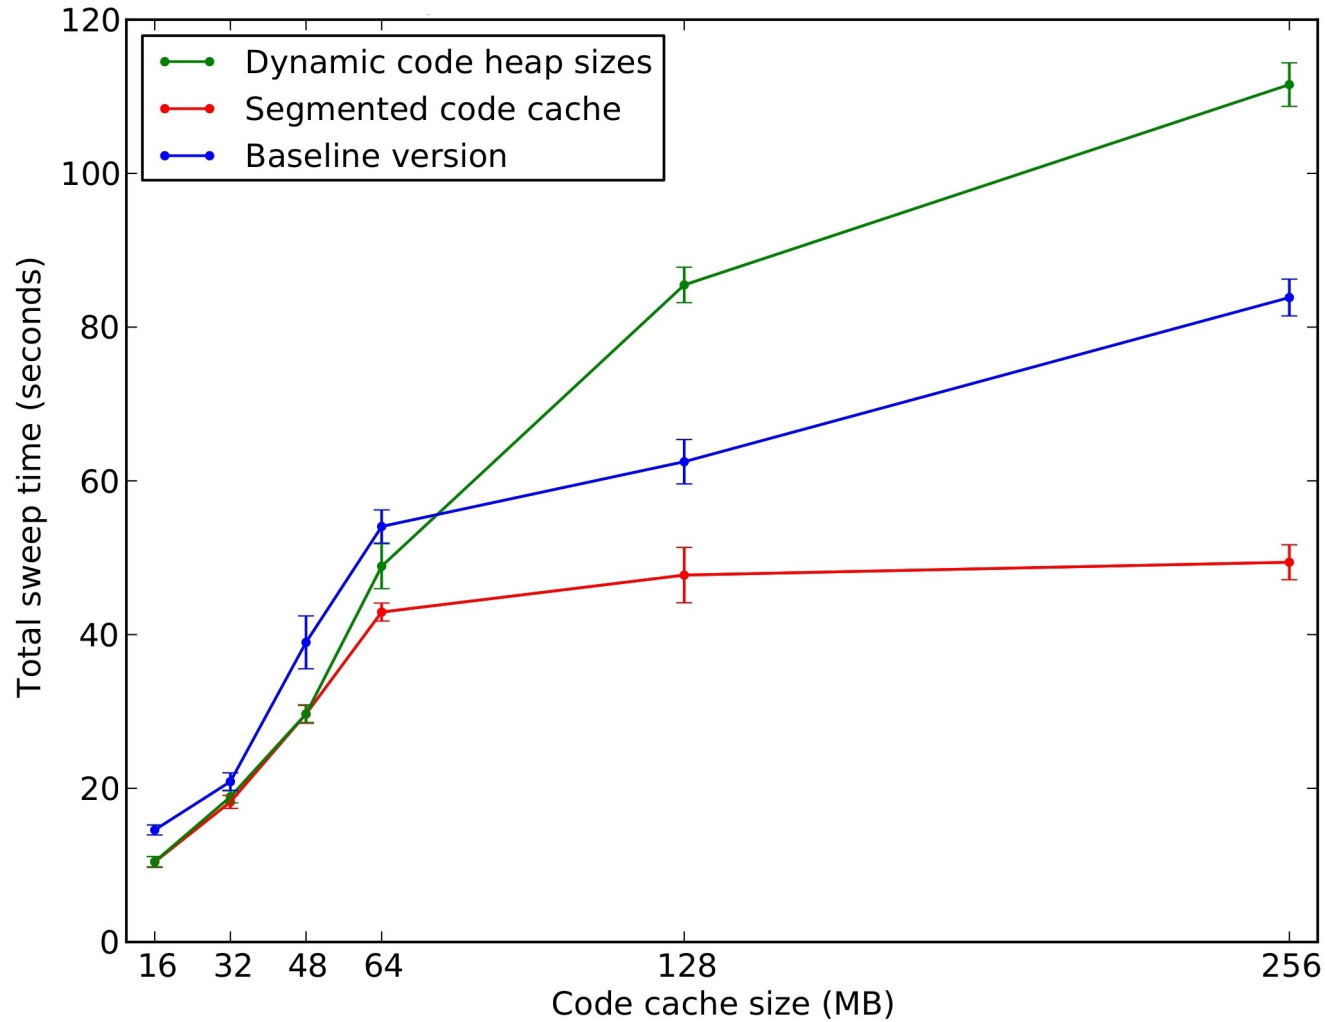


Instruction cache

- 14 %




Sweep time



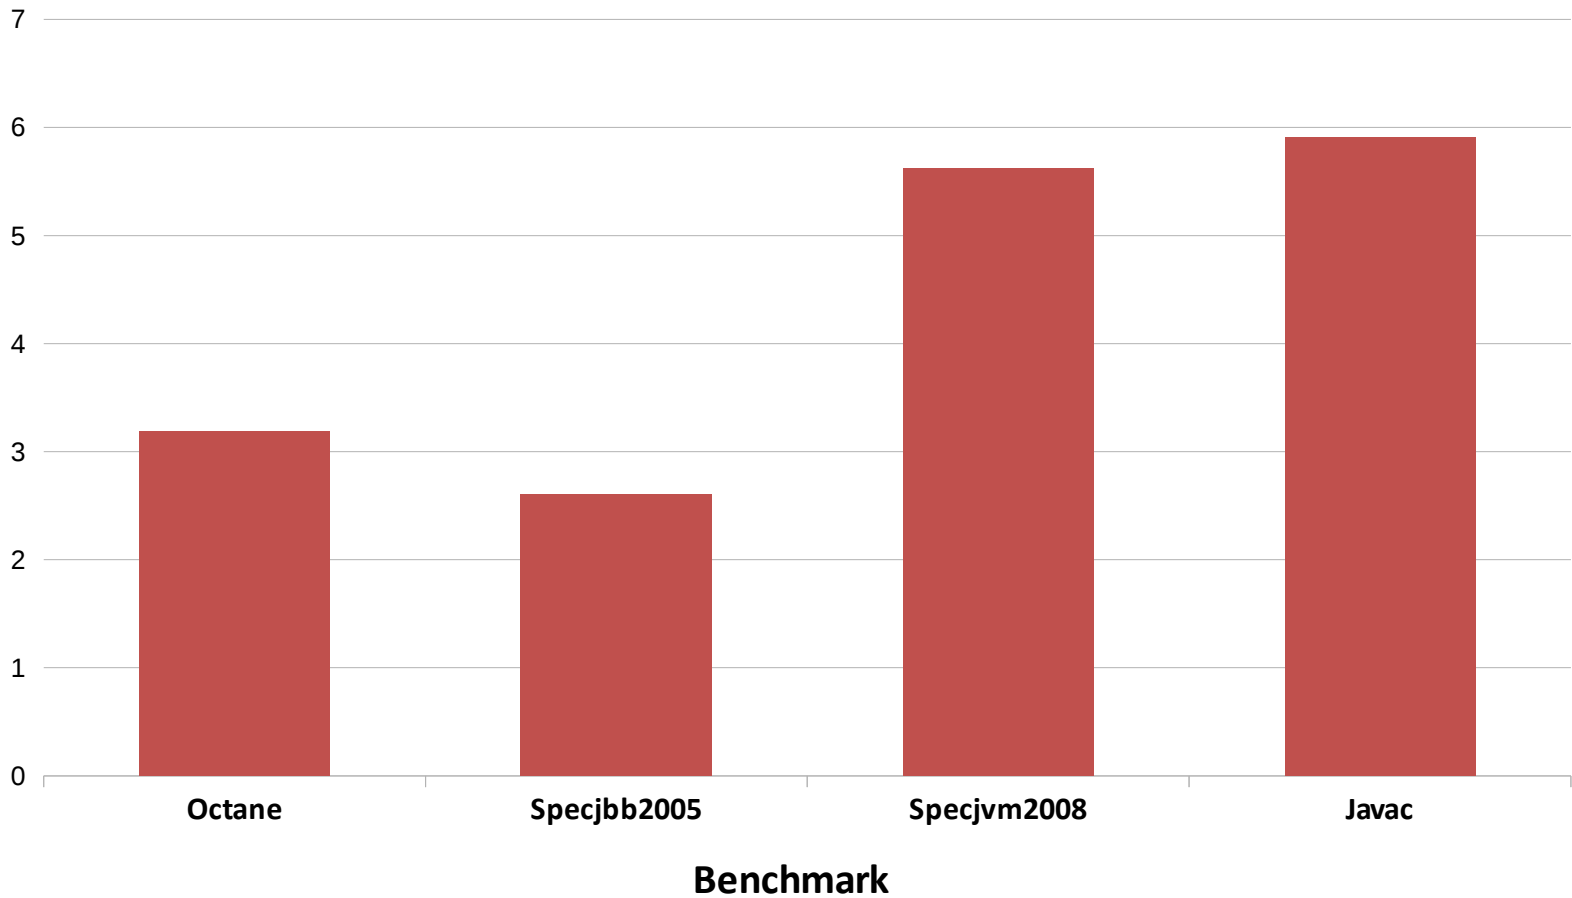
Sweep time

- 46 %



Execution time

Improvement in %



Evaluation summary

- **Performance **improvement** for regular sizes**

- Execution time: up to 6%
- Sweep time: up to 46%
- Fragmentation: up to 98%
- iTLB and iCache miss rates: up to 44%, 14%

- **Resizing does not pay off**

- **Only enable segmentation with**

- Tiered compilation
- Large code cache (> 240 MB)

Conclusion

- **Organization of code cache important**
 - Code locality
 - Fragmentation
- **Impact on overall performance**
- **Fully integrated into latest version**
 - Including tool support
 - Integration into JDK 9 in process

Future work

- **Separation** of code and metadata
- **Fine grained sweeping**
 - Sweep profiled code heap more often
- **Code heap partitioning**
- **Heterogeneous code**
 - More code heaps

Thank you for your attention!

<http://openjdk.java.net/jeps/197>

Related work

- **Java Virtual Machines**

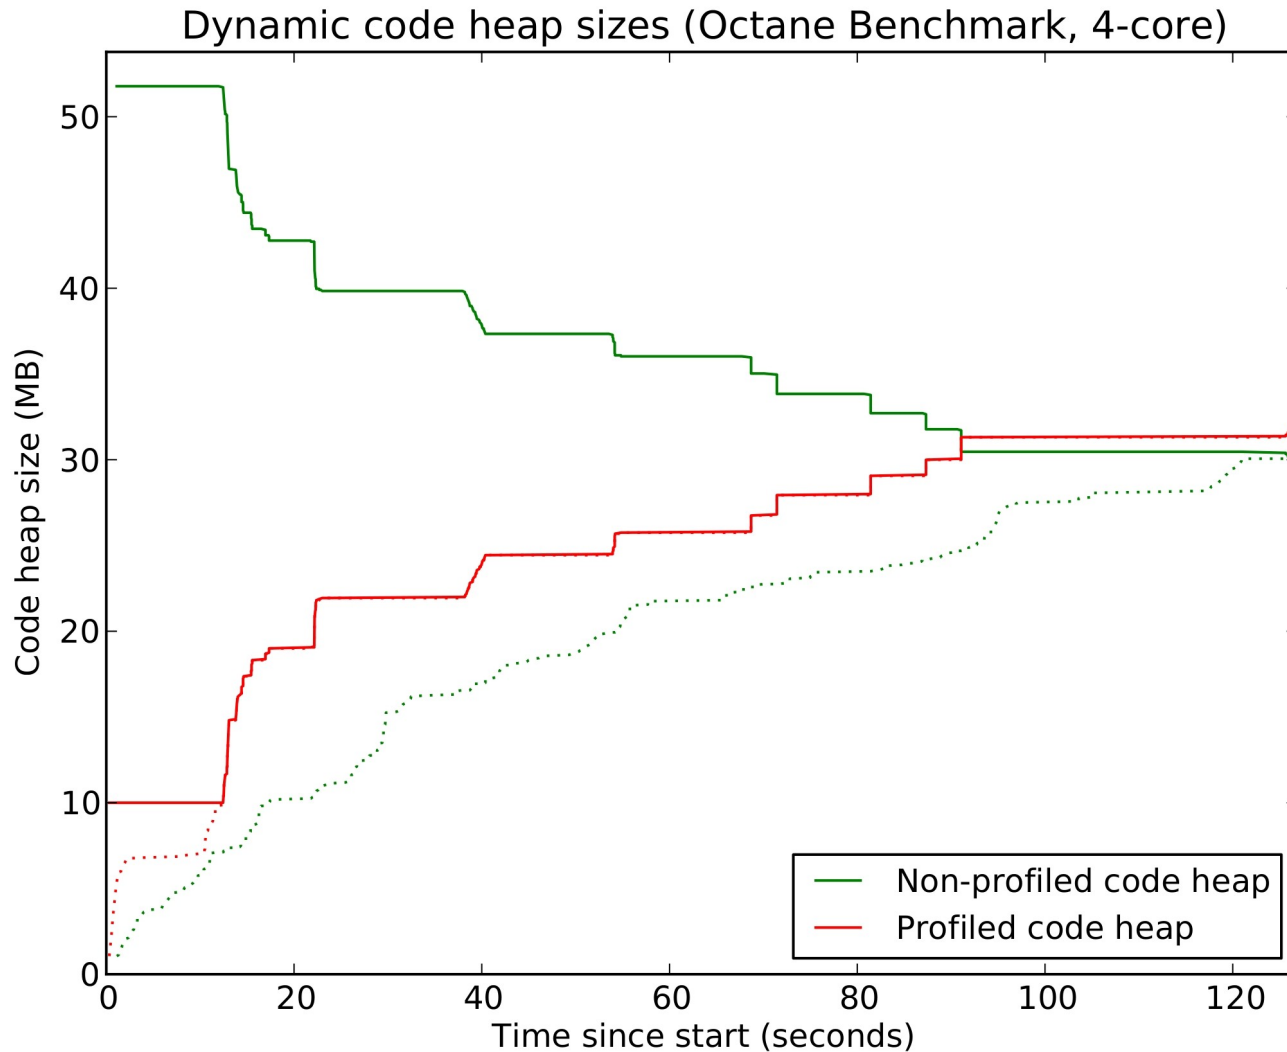
- Jikes RVM
- Maxine JVM
- Dalvik JVM

- **Dynamic Binary Translators**

- **Generational** code cache [Hazelwood and Smith]

- **Garbage collectors**

Resizing of code heaps



Resizing of code heaps

